

COMP219: Artificial Intelligence

Lecture 10: Heuristic Search

Overview


- Last time
 - Depth-limited, iterative deepening and bi-directional search
 - Avoiding repeated states
- Today
 - Show how applying **knowledge of the problem** can help
 - Introduce uniform cost search: dependent on the cost of each node
 - Introduce **heuristics**: rules of thumb
 - Introduce **heuristic search**
 - Greedy search
 - A* search
- Learning outcome covered today:
Identify, contrast and apply to simple examples the major search techniques that have been developed for problem-solving in AI

Real Life Problems


- Whatever search technique we use, we have **exponential** time complexity
- Tweaks to the algorithm will not reduce this to **polynomial**
- We need **problem specific** knowledge to **guide** the search
- Simplest form of problem specific knowledge is **heuristics**
- Usual implementation in search is via an **evaluation function** which indicates desirability of expanding a node

Path Cost Function

Recall: we have a path cost function, which gives the cost of each path. This is comprised of the step costs of each action on the path.



Finding The Best Paths



- Why not expand the *cheapest path first*?
- Intuition: cheapest is likely to be best!
- Performance is like breadth-first search but we use the minimum cost path rather than the shallowest to expand
- **Uniform cost search** orders the agenda as a priority queue using the lowest path cost of a node

Cheapest First

Choose Zerind

Choose Timisoara

Choose Sibiu

Choose Oradea

Choose Rimnicu Vilcea

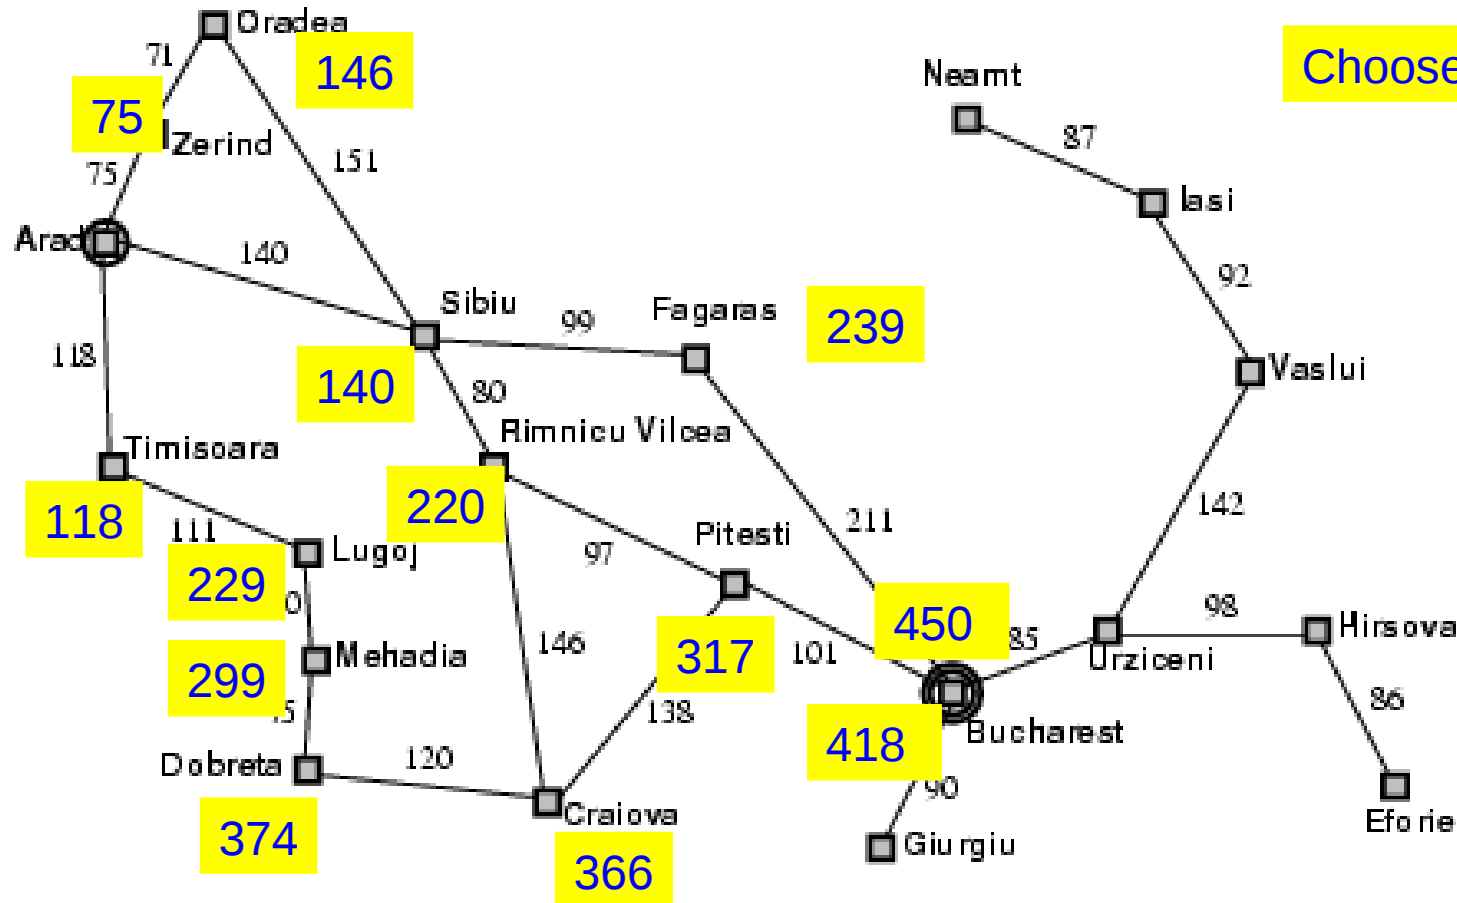
Choose Lugoj

Choose Fagaras

Choose Mehadia

Choose Pitesti

...



General Algorithm for Uniform Cost Search

```
agenda = initial state;  
while agenda not empty do  
    take node from agenda such that  
         $g(\text{node}) = \min \{ g(n) \mid n \text{ in agenda} \}$   
    if node is goal state then  
        return solution;  
    new nodes = apply operations to node;  
    add new nodes to the agenda;
```

Properties of Uniform Cost Search

- Uniform cost search **guaranteed** to find **cheapest** solution assuming path costs grow **monotonically**, i.e. the cost of a path never decreases as we move along it
- In other words, adding another step to the solution makes it more costly, i.e.
$$g(\text{successor}(n)) > g(n).$$
- If path costs don't grow monotonically, then exhaustive search is required
- Still requires many nodes to be examined

Informed Strategies



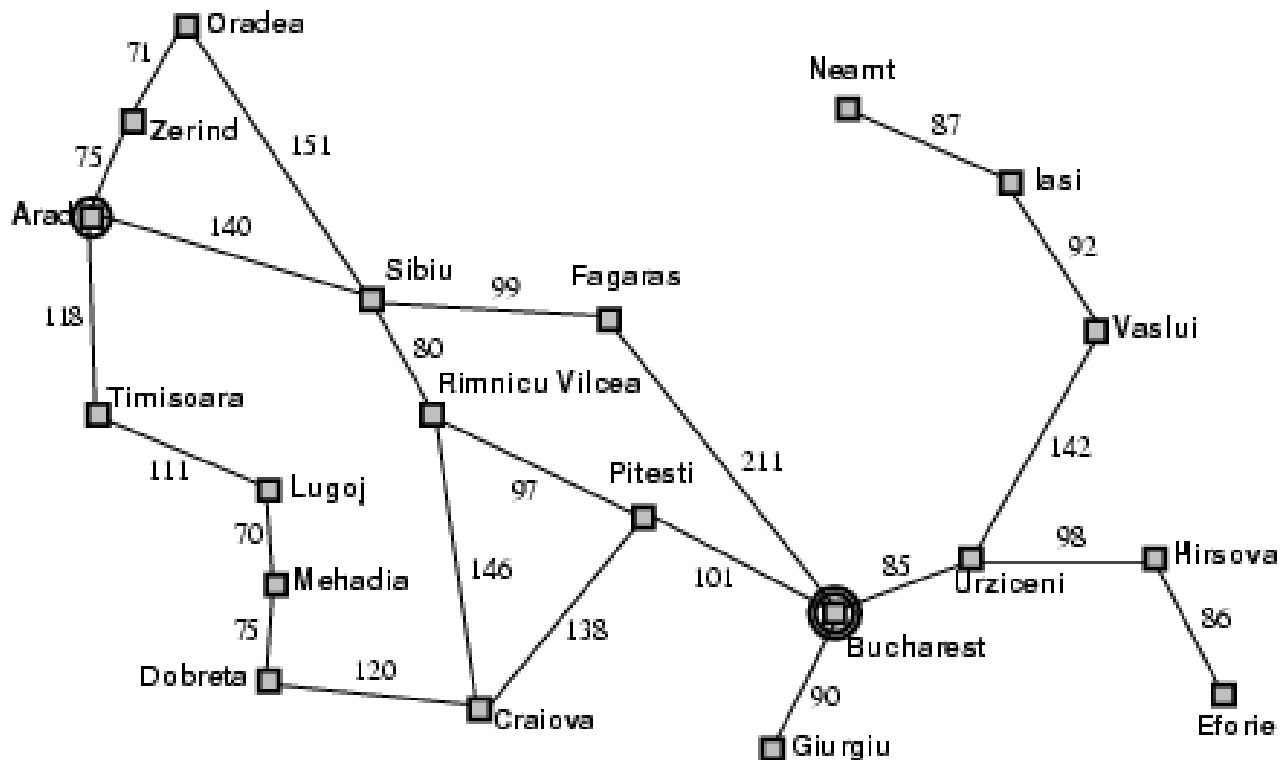
- Use problem-specific knowledge
- More efficient than blind search
- The most *promising* path first!
- Rather than trying all possible search paths, you try to focus on paths that seem to be getting you nearer your target/goal state



Greedy Search

- Most heuristics *estimate cost of cheapest path from node to solution*
- We have a *heuristic function*,
 $h : \text{Nodes} \rightarrow \mathbb{R}$
which estimates the distance from the node to the goal
- h can be any function but should have $h(n) = 0$ if n is a goal
- Example: In route finding, heuristic might be straight line distance from node to destination
- Greedy search expands the node that *appears to be* closest to goal

Romania Example



Straight-line distance
to Bucharest

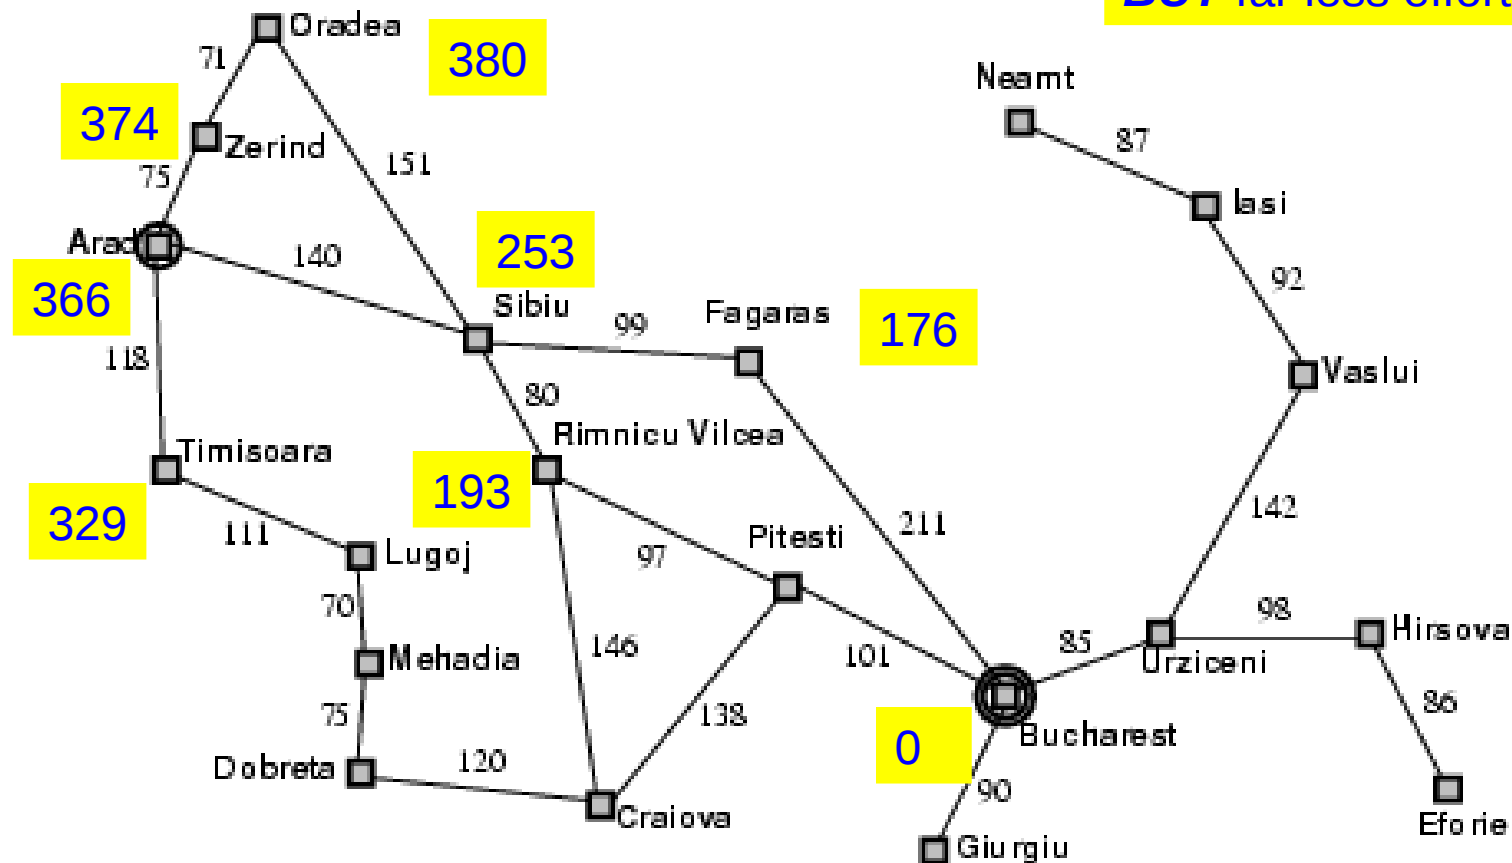
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy Search Example

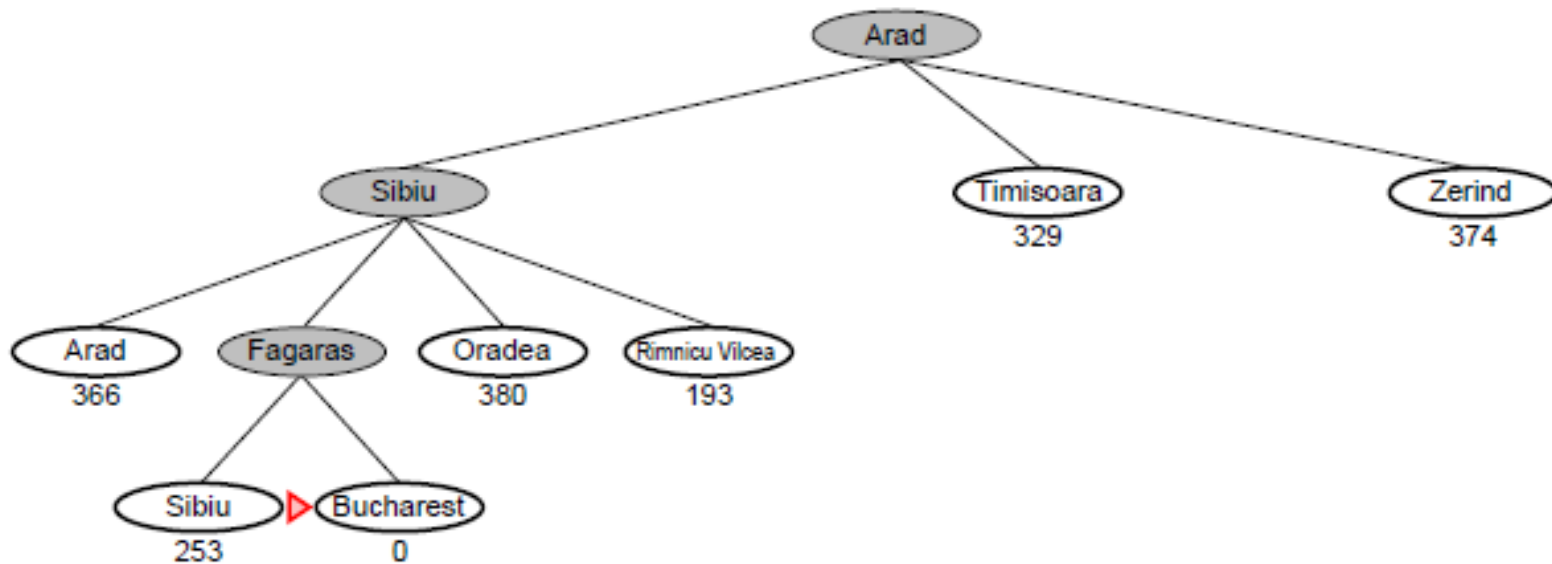
Choose Sibiu

Choose Fagaras

NOT the cheapest path
BUT far less effort



Search Tree



Exercise

General Algorithm for Greedy Search

```
agenda = initial state;  
while agenda not empty do  
    take node from agenda such that  
         $h(\text{node}) = \min \{ h(n) \mid n \text{ in agenda} \}$   
    if node is goal state then  
        return solution;  
    new nodes = apply operations to node;  
    add new nodes to the agenda;
```

Properties of Greedy Search

- Greedy search finds solutions quickly
- Doesn't always find the best
- May not find a solution if there is one (incomplete)
- Susceptible to false starts
- Only looking at *current node*. *Ignores past!*
- *Short sighted*



A* Search



- A* is a very efficient search strategy
- Basic idea is to **combine uniform cost** search and **greedy** search
- We look at the **cost so far** *and* the **estimated cost to goal**
- Gives heuristic **f**:

$$f(n) = g(n) + h(n)$$

- where
 - $g(n)$ is path cost of n
 - $h(n)$ is expected cost of cheapest solution from n
- Aims to **minimise overall cost**

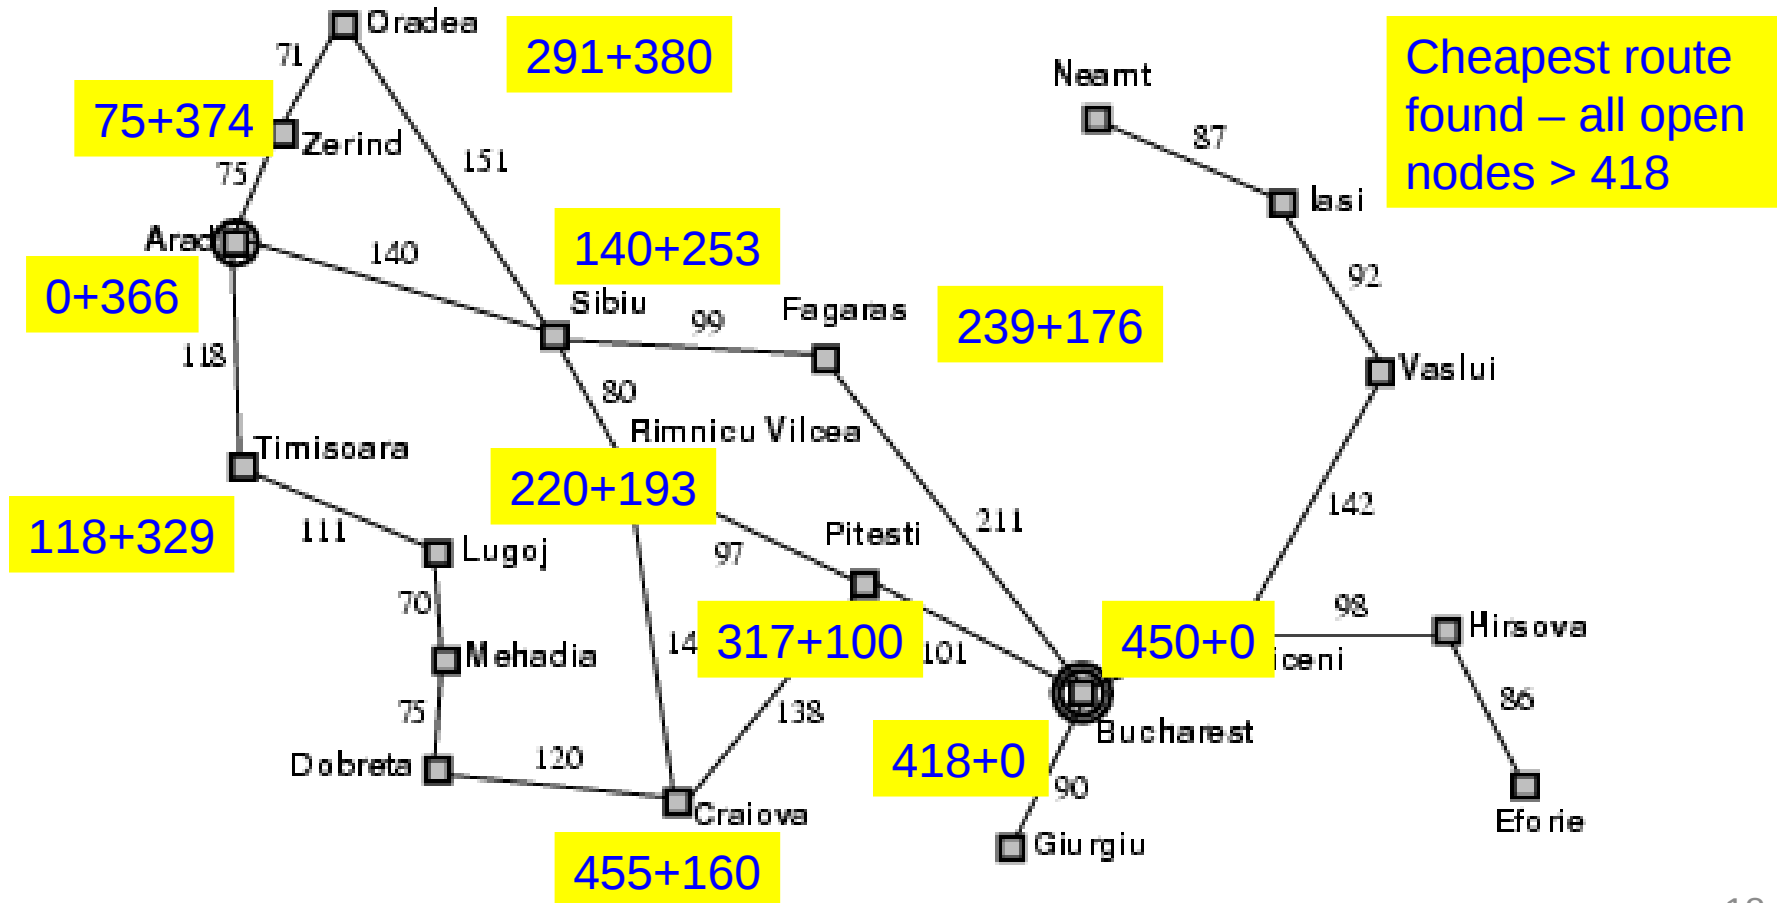
A* Search Example

Choose Sibiu

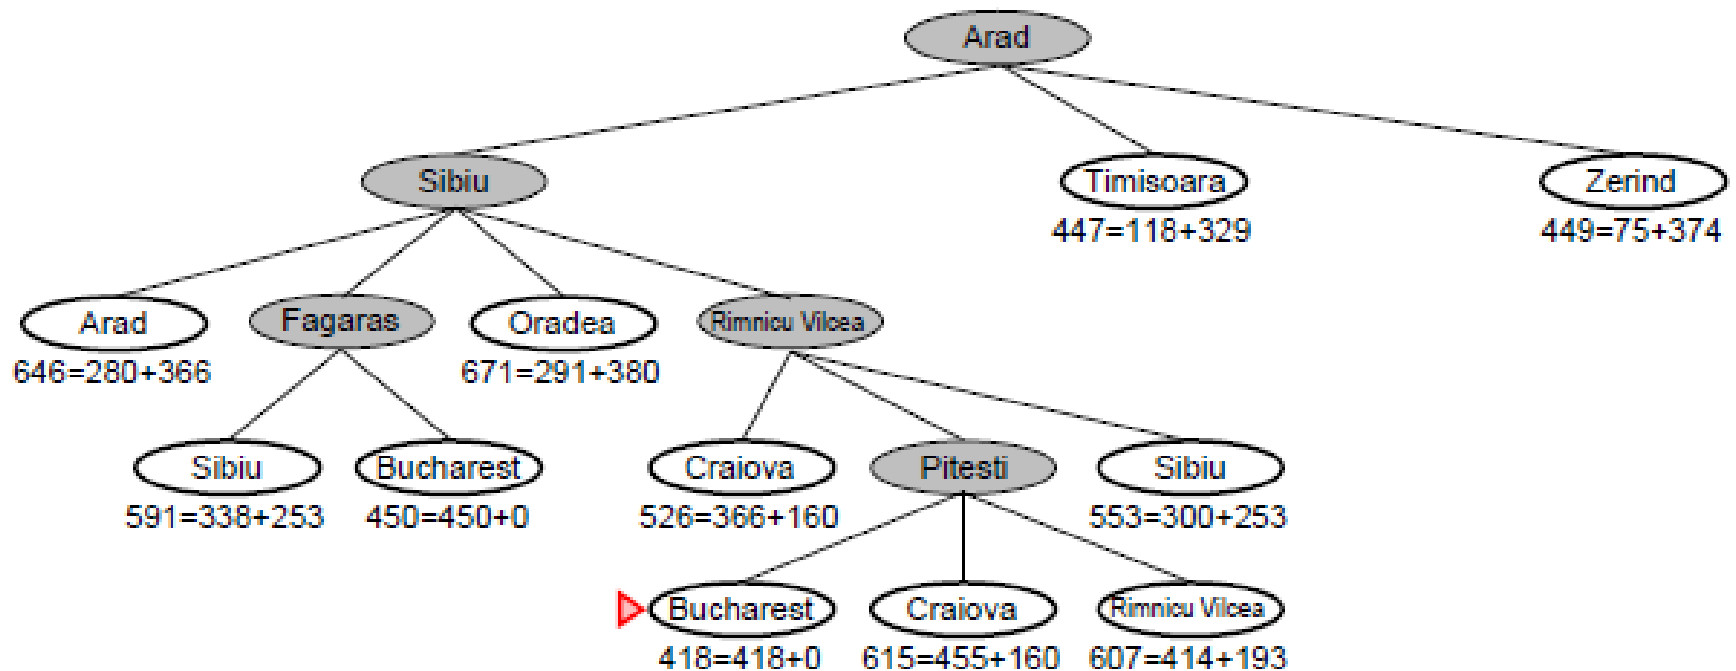
Choose Rimnicu Vilcea
 $413 < 415$

Choose Fagaras
 $415 < 417$

Choose Pitesti
 $418 < 450$



Search Tree



General Algorithm for A* Search

```
agenda = initial state;  
while agenda not empty do  
    take node from agenda such that  
         $f(\text{node}) = \min \{ f(n) \mid n \text{ in agenda} \}$   
        where  $f(n) = g(n) + h(n)$   
    if node is goal state then  
        return solution;  
    new nodes = apply operations to node;  
    add new nodes to the agenda;
```

Properties of A* Search

- Complete, provided
 - only **finitely** many nodes with $f < f(G)$
 - an **admissible** heuristic is used
 - **Never overestimates** the distance
 - i.e., $h(n) < \text{true}(n)$
where $\text{true}(n)$ is the true cost from n
 - Also require $h(n) \geq 0$, so $h(G) = 0$ for any goal G
- Exponential time
- Keeps all nodes in memory
- **Optimal**

Admissible Heuristics



- Example - for the 8-puzzle:

$h_1(n)$ = number of misplaced tiles

$h_2(n)$ = total *Manhattan distance*

- (i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

Exercise

Importance of the Heuristic Choice

- Typical search costs:
 - $d = 14$
 - IDS = 3,473,941 nodes
 - $A^*(h_1) = 539$ nodes
 - $A^*(h_2) = 113$ nodes
 - $d = 24$
 - IDS $\approx 54,000,000,000$ nodes
 - $A^*(h_1) = 39,135$ nodes
 - $A^*(h_2) = 1,641$ nodes

Summary

- Heuristic functions estimate costs of shortest paths
- Good heuristics can dramatically reduce search cost
- Greedy best-first search expands lowest h
 - incomplete and not always optimal
- A^* search expands lowest $g + h$
 - complete and optimal
 - also optimally efficient
- Next week
 - Search in complex environments (partial observation) and in game playing